

AOS2– Deep learning

Lecture 05: Attention Neural Networks

Sylvain Rousseau

What is attention

- Selectively concentrating on a few things
- Lead to breakthroughs on NLP, vision, speech recognition
 - First applied to neural networks in Bahdanau, Cho, and Bengio 2016 on top of RNN
 - Transformer architecture in Vaswani et al. 2017 and self-attention
 - In vision in Xu et al. 2016 with attention between images and captions
 - Vision transformer in Dosovitskiy et al. 2021

Attention

- Parametrized transform of a set of n vectors in \mathbb{R}^p into a set of m vectors in \mathbb{R}^p

$$U^T = [\mathbf{v}_1, \dots, \mathbf{v}_n] \xrightarrow{\Theta} Y^T = [\mathbf{y}_1, \dots, \mathbf{y}_m]$$

- Linear transform parametrized by **attention coefficients**

$$\mathbf{y}_j = \sum_{k=1}^n a_{jk} \mathbf{v}_k \quad \text{for } j = 1, \dots, m \quad \text{or} \quad Y = AU \quad \text{with} \quad A \in \mathbb{R}^{m \times n}$$

- Attention score a_{jk} reads: how \mathbf{v}_k is relevant for new vector \mathbf{y}_j
- Attention is usually calculated from a **softmax on attention scores**

$$a_{jk} = \frac{\exp(s_{jk})}{\sum_{l=1}^n \exp(s_{jl})} \quad A = \text{Softmax}^{\leftrightarrow}(S)$$

- Softmax on rows of score matrix $S \in \mathbb{R}^{m \times n}$

$$U^T = [\mathbf{v}_1, \dots, \mathbf{v}_n] \xrightarrow{\Theta = (\theta_1, \dots, \theta_m)} Y^T = [\mathbf{y}_1, \dots, \mathbf{y}_m]$$

- Contribution of \mathbf{v}_k in output \mathbf{y}_j depends on the content \mathbf{v}_k and a parameter θ_j

$$s_{jk} = s_{jk}(\mathbf{v}_k, \theta_j)$$

- For a given j , contribution of \mathbf{v}_k for \mathbf{y}_j depends only on the content \mathbf{v}_k

$$\mathcal{U}^T = [\mathbf{v}_1, \dots, \mathbf{v}_n] \xrightarrow{C^T = [\mathbf{c}_1, \dots, \mathbf{c}_m], \theta_c} \mathcal{Y}^T = [\mathbf{y}_1, \dots, \mathbf{y}_m]$$

- Scores depend on content \mathbf{v}_k , context \mathbf{c}_j and parameters θ_c

$$s_{jk} = s_{jk}(\mathbf{v}_k, \mathbf{c}_j, \theta_c)$$

- \mathbf{c}_j is data not a parameter!

Application to sequence to sequence modeling

- Machine translation

Translate a text from one language to another

- Video captioning

Generate a caption describing a video

- Open-ended question answering

Give a statement as an answer


- Summarization

Provide a summary of a long text

- Encoder–Decoder architecture


- Input sequence is entirely consumed

- All information is packed in last hidden state: bottleneck problem



Attention in Bahdanau, Cho, and Bengio 2016

- Context attention with $n = T, m = 1$
- Values are successive hidden states of RNN



Attention matrix

From english to french

- Give context to output correct token
- Attention matrix is aligning sequences




Figure 1: from Bahdanau, Cho, and Bengio 2016

Results and limitations

- BLEU score Models with attention have a better BLEU score
- Limitations
 - Still using RNN
 - One level attention only




Figure 2: from Bahdanau, Cho, and Bengio 2016

Transformer architecture

Transformer architecture

- Introduced in Vaswani et al. 2017
- Main features are:
 - Encoder-Decoder architecture
 - No RNN, key-value (self-)attention only
 - Multi-layered attention
 - Skip connections
 - Layer normalization
 - Positional encoding

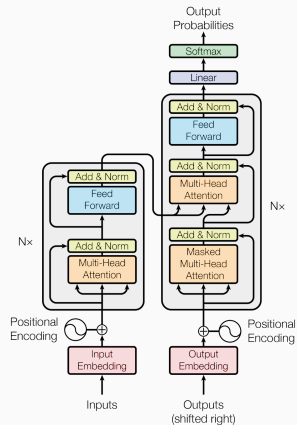


Figure 3: From Vaswani et al. 2017

Key-value attention

- Same as context attention except that:
 - Vector keys \mathbf{k}_k are used instead of \mathbf{v}_k when calculating the score
 - Context vectors that are now called query vectors \mathbf{q}_j

$$\mathcal{V}^T = [\mathbf{v}_1, \dots, \mathbf{v}_n] \xrightarrow[\mathcal{K}^T = [\mathbf{k}_1, \dots, \mathbf{k}_n], \theta_k]{\mathcal{Q}^T = [\mathbf{q}_1, \dots, \mathbf{q}_m], \theta_q} \mathcal{Y}^T = [\mathbf{y}_1, \dots, \mathbf{y}_m]$$

- Attention score is then

$$s_{jk} = s_{jk}(\mathbf{k}_k, \mathbf{q}_j, \theta)$$

- Differentiable database

Attention scores

- Additive attention (basically a 2-layers MLP in Bahdanau, Cho, and Bengio 2016)

$$a_{jk} = \mathbf{w}_a \cdot \tanh(\mathbf{w}_q \mathbf{q}_j + \mathbf{w}_k \mathbf{k}_k)$$

- Bilinear attention in Luong, Pham, and Manning 2015

$$s_{jk} = \mathbf{q}_j \mathbf{W} \mathbf{k}_k$$

- Dot-product attention score Luong, Pham, and Manning 2015

$$s_{jk} = \mathbf{q}_j \cdot \mathbf{k}_k$$

- Scaled dot-product attention score in Vaswani et al. 2017

$$s_{jk} = \frac{\mathbf{q}_j \cdot \mathbf{k}_k}{\sqrt{d}}$$

Key-value attention: example

- Values

$$\mathbf{v}_1 = \begin{pmatrix} 2 \\ 3 \\ 1 \end{pmatrix}, \quad \mathbf{v}_2 = \begin{pmatrix} 2 \\ -1 \\ 0 \end{pmatrix}, \quad \mathbf{v}_3 = \begin{pmatrix} 0 \\ 5 \\ 1 \end{pmatrix} \quad \text{so that} \quad \mathbf{V} = \begin{pmatrix} 2 & 3 & 1 \\ 2 & -1 & 0 \\ 0 & 5 & 1 \end{pmatrix},$$

- Keys

$$\mathbf{k}_1 = \begin{pmatrix} 2 \\ 1 \\ -1 \end{pmatrix}, \quad \mathbf{k}_2 = \begin{pmatrix} 0 \\ 3 \\ -1 \end{pmatrix}, \quad \mathbf{k}_3 = \begin{pmatrix} 1 \\ 1 \\ 3 \end{pmatrix} \quad \text{so that} \quad \mathbf{K} = \begin{pmatrix} 2 & 1 & -1 \\ 0 & 3 & -1 \\ 1 & 1 & 3 \end{pmatrix},$$

- Same number of keys and values
- Not necessarily same dimension

Key-value attention: example

- Query

$$\mathbf{q}_1 = \begin{pmatrix} 2 \\ -1 \\ 0 \end{pmatrix}, \quad \mathbf{q}_2 = \begin{pmatrix} -2 \\ 1 \\ 4 \end{pmatrix} \quad \text{so that} \quad Q = \begin{pmatrix} 2 & -1 & 0 \\ -2 & 1 & 4 \end{pmatrix},$$

- Any number of query but same dimension as that of keys
- Dot-product attention: $s_{jk} = \langle \mathbf{k}_k, \mathbf{q}_j \rangle$ so that the score matrix is $S = QK^T$

$$Q = \begin{pmatrix} 2 & -1 & 0 \\ -2 & 1 & 4 \end{pmatrix}, \quad K = \begin{pmatrix} 2 & 1 & -1 \\ 0 & 3 & -1 \\ 1 & 1 & 3 \end{pmatrix}, \quad \text{so that} \quad QK^T = \begin{pmatrix} 3 & -3 & 1 \\ -7 & -1 & 11 \end{pmatrix},$$

Key-value attention: example

- Softmax on rows of $S = QK^T$

$$A = \begin{pmatrix} 0.879 & 0.002 & 0.119 \\ 0.0 & 0.0 & 1.0 \end{pmatrix}$$

- From the attention matrix A , \mathbf{v}_1 and \mathbf{v}_3 are approximately selected

$$X' = AV = \begin{pmatrix} 1.762 & 3.23 & 0.998 \\ 0.0 & 5.0 & 1.0 \end{pmatrix}, \quad V = \begin{pmatrix} 2 & 3 & 1 \\ 2 & -1 & 0 \\ 0 & 5 & 1 \end{pmatrix}$$

Self-attention

How to calculate context or key and query vectors?

Answer: we use the content itself!

- Context-based self attention: context is calculated from \mathbf{v}_j itself!

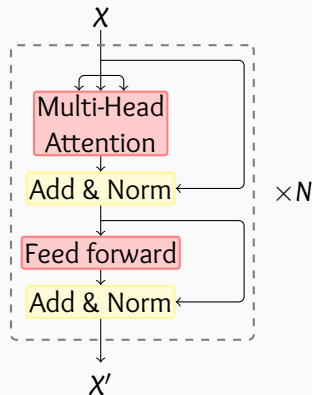
$$\mathbf{c}_j = f(\mathbf{v}_j, \theta_j)$$

- Key-value self attention: both key and query vectors are calculated from \mathbf{v}_j itself!

$$\mathbf{k}_j = f^k(\mathbf{v}_j, \theta_j^k) \quad \mathbf{q}_j = f^q(\mathbf{v}_j, \theta_j^q)$$

Transformer encoder block

- Main features
 - Key-value self attention
 - Layer normalization
 - Multi-head attention
 - Skip connections
 - Transformer blocks are stacked



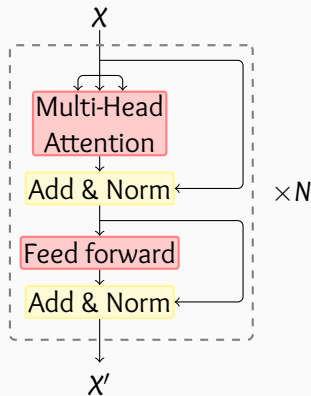
Key-value self-attention

- $X \in \mathbb{R}^{n \times p}$ contains a sequence of n tokens embedded in \mathbb{R}^p

$$X^T = [\mathbf{x}_1, \dots, \mathbf{x}_n]$$

- Keys, values and queries are created from X ,
 - $V = XW_V$ with $W_V \in \mathbb{R}^{p \times d_v}$
 - $Q = XW_Q$ with $W_Q \in \mathbb{R}^{p \times d_k}$
 - $K = XW_K$ with $W_K \in \mathbb{R}^{p \times d_k}$
- Simple key-value self-attention

$$T = \text{Attention}(V, K, Q) = \text{Softmax}^{\leftrightarrow} \left(\frac{QK^T}{\sqrt{d_k}} \right) \cdot V$$



Layer normalization from Ba, Kiros, and Hinton 2016

- Standardizing each token representation

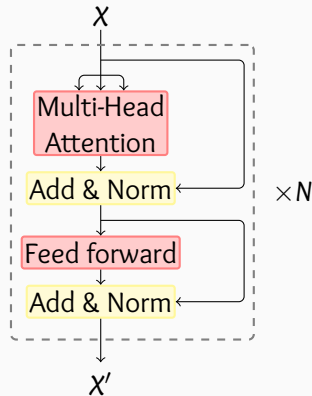
$$\text{LayerNorm}(\mathbf{x}_t) = \gamma \odot \frac{\mathbf{x}_t - \mu_t}{\sigma_t} + \mathbf{b}$$

- Rescaling parameters $\gamma, \mathbf{b} \in \mathbb{R}^p$
- μ_t the sample mean of \mathbf{x}_t :

$$\mu_t = \frac{1}{p} \sum_{i=0}^{p-1} (\mathbf{x}_t)_i$$

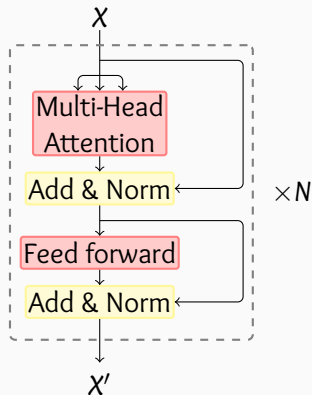
- σ_t the sample standard deviation of \mathbf{x} :

$$\sigma_t = \sqrt{\frac{1}{p} \sum_{i=0}^{p-1} ((\mathbf{x}_t)_i - \mu_t)^2}$$



Feed forward

- Position-wise 2-layers FFN
- Same parameters for each position
- Different parameters between transformer encoder block



Summary: single head transformer block

A transformer block is a map $X \in \mathbb{R}^{n \times p} \mapsto X' \in \mathbb{R}^{n \times p}$

- Keys, values and queries are created from X , $W_Q, W_K \in \mathbb{R}^{p \times d_k}, W_V \in \mathbb{R}^{p \times d_v}$
 $Q = XW_Q \in \mathbb{R}^{n \times d_k}, \quad K = XW_K \in \mathbb{R}^{n \times d_k}, \quad V = XW_V \in \mathbb{R}^{n \times d_v}$

- Attention matrix

$$T = \text{Attention}(Q, K, V) = \text{Softmax}^{\leftrightarrow} \left(\frac{QK^T}{\sqrt{d_k}} \right) \cdot V \in \mathbb{R}^{n \times d_v}$$

- Residual mapping and layer normalization, $W_0 \in \mathbb{R}^{d_v \times p}$
 $U = \text{LayerNorm}(TW_0 + X) \in \mathbb{R}^{n \times p}$
- Position-wise FFN, $W_1 \in \mathbb{R}^{p \times d_{\text{ff}}}, W_2 \in \mathbb{R}^{d_{\text{ff}} \times p}$

$$Z = \text{ReLU}(UW_1)W_2 \in \mathbb{R}^{n \times p}$$

- Residual mapping and layer normalization

$$X' = \text{LayerNorm}(Z + U) \in \mathbb{R}^{n \times p}$$

Multi-head transformer block

- X is first split into H pieces

$$X = [X_1, \dots, X_H], \quad X_i \in \mathbb{R}^{n \times p_h} \text{ with } H \cdot p_h = p$$

- Keys, values and queries are created from each X_h , $w_Q^{(h)}, w_K^{(h)} \in \mathbb{R}^{p \times d_{k_h}}, w_V^{(h)} \in \mathbb{R}^{p \times d_{v_h}}$

$$Q^{(h)} = X_h w_Q^{(h)} \in \mathbb{R}^{n \times d_{k_h}}, \quad K^{(h)} = X_h w_K^{(h)} \in \mathbb{R}^{n \times d_{k_h}}, \quad V^{(h)} = X_h w_V^{(h)} \in \mathbb{R}^{n \times d_{v_h}}$$

- Multi-head attention matrix

$$T^{(h)} = \text{Attention} \left(Q^{(h)}, K^{(h)}, V^{(h)} \right) = \text{Softmax}^{\leftrightarrow} \left(\frac{Q^{(h)} (K^{(h)})^T}{\sqrt{d_{k_h}}} \right) \cdot V^{(h)} \in \mathbb{R}^{n \times d_{v_h}}$$

$$T = [T^{(1)}, \dots, T^{(H)}] \in \mathbb{R}^{n \times d_v} \quad \text{with } d_v = H \cdot d_{v_h}$$

Positional encoding

- Equivariant to permutation of samples:
 - If X gives Y , $X_\sigma = P_\sigma X$ gives $Y_\sigma = P_\sigma Y$
 - We can check that $X'_\sigma = P_\sigma X'$
- Temporal information is not taken into account
- Two different strategies:
 - Learnable positional encoding
 - Enlarge embedding with parameters tied to position
 - Problem at test time with unseen positions during train time
 - Non-learnable positional encoding
 - No extra parameters
 - Defined for all length of sequence even if it has not been seen in the train set.

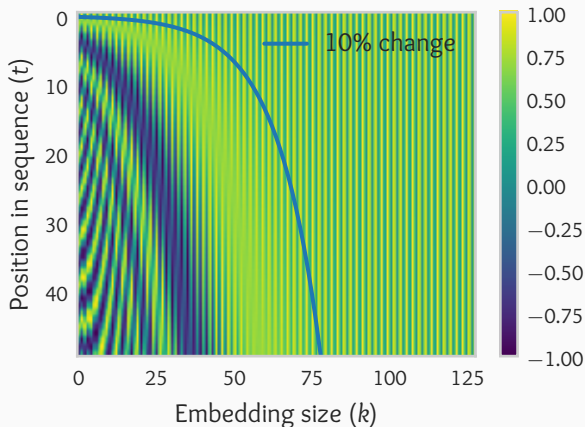
Positional encoding

- $\mathbf{e}_1, \dots, \mathbf{e}_n$ is a sequence of vectors in \mathbb{R}^d
- positional encoding $\mathbf{p}_t \in \mathbb{R}^d$ of \mathbf{e}_t is

$$p_t^k = \begin{cases} \sin 2\pi t / T_{2i} & \text{if } k = 2i \\ \cos 2\pi t / T_{2i} & \text{if } k = 2i + 1 \end{cases}$$


where $T_k = 2\pi \cdot 10000^{k/d}$

- Input after positional encoding is $\mathbf{e}_t + \mathbf{p}_t$ instead of \mathbf{e}_t




Translation invariance

- $\langle p_t, p_{t'} \rangle$ only depends on $|t - t'|$
- $$\begin{aligned}\langle p_t, p_{t'} \rangle &= \sum_{k \text{ even}} \sin\left(\frac{2\pi t}{T_k}\right) \sin\left(\frac{2\pi t'}{T_k}\right) \\ &\quad + \cos\left(\frac{2\pi t}{T_k}\right) \cos\left(\frac{2\pi t'}{T_k}\right) \\ &= \sum_{k \text{ even}} \cos\left(\frac{2\pi(t - t')}{T_k}\right) \\ &= \sum_{k \text{ even}} \cos\left(\frac{2\pi |t - t'|}{T_k}\right)\end{aligned}$$



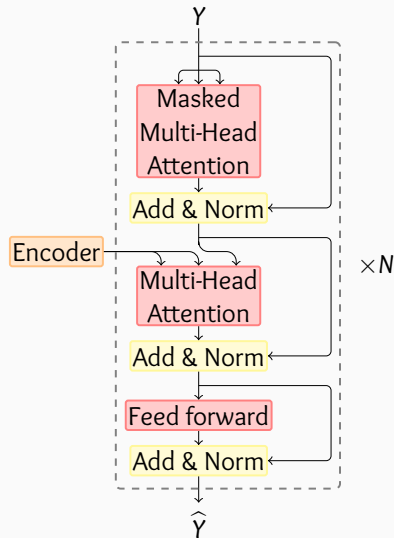
Variance of positional encoding

- Plot of $\text{Var}_{|t-t'| \sim \mathcal{U}(50)} \langle \mathbf{p}_t^{1\dots k}, \mathbf{p}_{t'}^{1\dots k} \rangle$ w.r.t k
- Position information is encoded in first $\simeq 50$ elements with $d = 128$



Transformer decoder block

- Main features
 - Masked key-value self attention
 - Layer normalization
 - Skip connections
 - Transformer blocks are stacked
- Two main differences
 - Attention matrix is masked out to prevent receiving attention from future
 - Keys and values come from the encoder and decoder make queries on these



Masked multi-head attention

- What is the loss governing the whole architecture?

| | | | | | | |
|------------|------------------------------|----------------------|----------------------|---------|--------------------------|------------------------------|
| Source | $\langle \text{bos} \rangle$ | \mathbf{y}_1 | \mathbf{y}_2 | \dots | \mathbf{y}_{m-2} | \mathbf{y}_{m-1} |
| | \downarrow | \downarrow | \downarrow | \dots | \downarrow | \downarrow |
| Prediction | $\hat{\mathbf{y}}_1$ | $\hat{\mathbf{y}}_2$ | $\hat{\mathbf{y}}_3$ | \dots | $\hat{\mathbf{y}}_{m-1}$ | $\hat{\mathbf{y}}_m$ |
| Target | \mathbf{y}_1 | \mathbf{y}_2 | \mathbf{y}_3 | \dots | \mathbf{y}_{m-1} | $\langle \text{eos} \rangle$ |

- Loss on couples $(\hat{\mathbf{y}}_1, \mathbf{y}_1), \dots, (\hat{\mathbf{y}}_{m-1}, \mathbf{y}_{m-1})$ and $(\hat{\mathbf{y}}_m, \langle \text{eos} \rangle)$
- But $\hat{\mathbf{y}}_1$ should not be able to receive attention from \mathbf{y}_2 !
- Masked attention prevents tokens from receiving attention from future ones

Masked multi-head attention

- How to prevent tokens from receiving attention from future ones?
- Change attention scores $S = QK^T$ so that Softmax ignore them
 - $\text{Softmax}(2, 2, -\infty) = (1/2, 1/2, 0)$
 - $\text{Softmax}(1, 1, 1, -\infty) = (1/3, 1/3, 1/3, 0)$

Single-head decoder block I

A decoder transformer block is a map $Y \in \mathbb{R}^{m \times p} \mapsto \hat{Y} \in \mathbb{R}^{m \times p}$

- Keys, values and queries are created from Y , $W_{V_1} \in \mathbb{R}^{p \times d_v}$, $W_{Q_1}, W_{K_1} \in \mathbb{R}^{p \times d_k}$

$$Q_1 = YW_{Q_1} \in \mathbb{R}^{m \times d_k}, \quad K_1 = YW_{K_1} \in \mathbb{R}^{m \times d_k}, \quad V_1 = YW_{V_1} \in \mathbb{R}^{m \times d_v}$$

- Masked attention matrix

$$T_1 = \text{Attention}(Q_1, K_1, V_1) = \text{Softmax}^{\leftrightarrow} \left(\frac{Q_1 K_1^T + M}{\sqrt{d_k}} \right) \cdot V_1 \in \mathbb{R}^{m \times d_v}$$

- Residual mapping and layer normalization, $W_0 \in \mathbb{R}^{d_v \times p}$

$$U_1 = \text{LayerNorm}(T_1 W_0 + Y) \in \mathbb{R}^{m \times p}$$

- Query of decoder on keys and values of the encoder

$$Q_2 = U_1 W_{Q_2} \in \mathbb{R}^{m \times d_k}, \quad K_2 = X' W_{K_2} \in \mathbb{R}^{n \times d_k}, \quad V_2 = X' W_{V_2} \in \mathbb{R}^{n \times d_v}$$

- Cross-attention with attention matrix $A \in \mathbb{R}^{m \times n}$

$$T_2 = \text{Attention}(Q_2, K_2, V_2) = \text{Softmax}^{\leftrightarrow} \left(\frac{Q_2 K_2^T}{\sqrt{d_k}} \right) \cdot V_2 \in \mathbb{R}^{m \times d_v}$$

Single-head decoder block II

- Residual mapping and layer normalization, $\mathcal{W}_1 \in \mathbb{R}^{d_v \times p}$

$$U_2 = \text{LayerNorm} (T_2 \mathcal{W}_1 + U_1) \in \mathbb{R}^{m \times p}$$

- Position-wise FFN, $\mathcal{W}_2 \in \mathbb{R}^{p \times d_{\text{ff}}}$, $\mathcal{W}_3 \in \mathbb{R}^{d_{\text{ff}} \times p}$

$$Z = \text{ReLU} (U_2 \mathcal{W}_2) \mathcal{W}_3 \in \mathbb{R}^{m \times p}$$

- Residual mapping and layer normalization,

$$\hat{Y} = \text{LayerNorm} (Z + U_2) \in \mathbb{R}^{m \times p}$$

- Set of m distributions on tokens, $\mathcal{W} \in \mathbb{R}^{p \times N}$

$$\text{Softmax}^{\leftrightarrow} (\hat{Y} \mathcal{W})$$

- [1] Minh-Thang Luong, Hieu Pham, and Christopher D. Manning. **“Effective approaches to attention-based neural machine translation”**. 2015. arXiv: 1508.04025.
- [2] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. **“Layer normalization”**. 2016. arXiv: 1607.06450.
- [3] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. **“Neural Machine Translation by Jointly Learning to Align and Translate”**. May 19, 2016. arXiv: 1409.0473 [cs, stat]. URL: <http://arxiv.org/abs/1409.0473> (visited on 10/14/2021).
- [4] Kelvin Xu et al. **“Show, Attend and Tell: Neural Image Caption Generation with Visual Attention”**. Apr. 19, 2016. arXiv: 1502.03044 [cs]. URL: <http://arxiv.org/abs/1502.03044> (visited on 12/15/2021).

- [5] Ashish Vaswani et al. **“Attention is all you need”**. In: *Advances in neural information processing systems*. 2017, pp. 5998–6008. arXiv: 1706.03762.
- [6] Alexey Dosovitskiy et al. **“An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale”**. June 3, 2021. arXiv: 2010.11929 [cs]. URL: <http://arxiv.org/abs/2010.11929> (visited on 12/15/2021).